

on-modeling-for-used-cars-in-india

November 26, 2024

[]:

1 Project Title: Price prediction modeling for used cars in India

1.0.1 Overview

This dataset contains information about used cars in the Indian market, comprising 9,582 entries with 11 detailed attributes. The data appears to be collected up to November 2024, providing a comprehensive view of the second-hand car market in India.

1.0.2 Dataset Features

- Brand: Car manufacturer (e.g., Volkswagen, Maruti Suzuki, Honda, Tata)
- Model: Specific car model (e.g., Taigun, Baleno, Polo, WRV)
- Year: Manufacturing year of the vehicle (ranging from older models to 2024)
- Age: Age of the vehicle in years
- kmDriven: Total kilometers driven by the vehicle
- Transmission: Type of transmission (Manual or Automatic)
- Owner: Ownership status (first or second owner)
- FuelType: Type of fuel (Petrol, Diesel, Hybrid/CNG)
- PostedDate: When the car listing was posted
- AdditionalInfo: Extra details about the vehicle
- AskPrice: Listed price in Indian Rupees ()

[]:

[]:

```
[6]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[149]: # loading dataset
```

```
data = pd.read_csv("used_car_dataset.csv")
```

```
[150]: data.head()
```

```
[150]:      Brand      model  Year  Age      kmDriven Transmission  Owner  \
0      Honda      City  2001   23    98,000 km      Manual  second
1      Toyota    Innova  2009   15  190000.0 km      Manual  second
2  Volkswagen  VentoTest  2010   14    77,246 km      Manual  first
3  Maruti Suzuki    Swift  2017    7    83,500 km      Manual  second
4  Maruti Suzuki   Baleno  2019    5    45,000 km  Automatic  first

      FuelType PostedDate      AdditionInfo  \
0    Petrol    Nov-24  Honda City v teck in mint condition, valid gen...
1    Diesel    Jul-24  Toyota Innova 2.5 G (Diesel) 7 Seater, 2009, D...
2    Diesel    Nov-24  Volkswagen Vento 2010-2013 Diesel Breeze, 2010...
3    Diesel    Nov-24      Maruti Suzuki Swift 2017 Diesel Good Condition
4    Petrol    Nov-24      Maruti Suzuki Baleno Alpha CVT, 2019, Petrol

      AskPrice
0    1,95,000
1    3,75,000
2    1,84,999
3    5,65,000
4    6,85,000
```

```
[151]: data.tail()
```

```
[151]:      Brand      model  Year  Age      kmDriven Transmission  Owner  \
9577   Skoda    Octavia  2014   10  105,904 km      Automatic  second
9578  Maruti Suzuki  Alto-800  2020    4    55,000 km      Manual  first
9579  Maruti Suzuki    Ritz  2013   11    92,000 km      Manual  first
9580   Hyundai    Verna  2019    5    72,000 km      Automatic  first
9581   Hyundai   New i20  2021    3    83,228 km      Manual  second

      FuelType PostedDate  \
9577    Diesel    Oct-24
9578 Hybrid/CNG    Nov-24
9579    Diesel    Nov-24
9580    Petrol    Oct-24
9581    Petrol    Nov-24

      AdditionInfo      AskPrice
9577  Skoda Octavia 1.9 Elegance TDI, 2014, Diesel    10,40,000
9578  Maruti Suzuki Alto 800 CNG LXI Optional, 2020,...    3,75,000
9579      Maruti Suzuki Ritz VDi, 2013, Diesel    4,15,000
9580  Hyundai Verna VTVT 1.6 AT SX Option, 2019, Petrol    8,55,000
9581      Hyundai New i20 1.2 Asta IVT, 2021, Petrol    6,99,000
```

```
[152]: data.shape
```

```
[152]: (9582, 11)
```

```
[153]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9582 entries, 0 to 9581
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Brand            9582 non-null   object
1   model            9582 non-null   object
2   Year             9582 non-null   int64
3   Age              9582 non-null   int64
4   kmDriven         9535 non-null   object
5   Transmission     9582 non-null   object
6   Owner            9582 non-null   object
7   FuelType         9582 non-null   object
8   PostedDate       9582 non-null   object
9   AdditionInfo     9582 non-null   object
10  AskPrice         9582 non-null   object
dtypes: int64(2), object(9)
memory usage: 823.6+ KB
```

```
[154]: # find the null values
```

```
data.isna().sum()
```

```
[154]: Brand            0
      model            0
      Year             0
      Age              0
      kmDriven         47
      Transmission     0
      Owner            0
      FuelType         0
      PostedDate       0
      AdditionInfo     0
      AskPrice         0
      dtype: int64
```

```
[155]: data.columns
```

```
[155]: Index(['Brand', 'model', 'Year', 'Age', 'kmDriven', 'Transmission', 'Owner',
        'FuelType', 'PostedDate', 'AdditionInfo', 'AskPrice'],
        dtype='object')
```

```
[156]: # drop the column

data.drop(columns = ['AdditionInfo'], axis=1, inplace=True)
```

```
[157]: data.head()
```

```
[157]:
```

	Brand	model	Year	Age	kmDriven	Transmission	Owner	\
0	Honda	City	2001	23	98,000 km	Manual	second	
1	Toyota	Innova	2009	15	190000.0 km	Manual	second	
2	Volkswagen	VentoTest	2010	14	77,246 km	Manual	first	
3	Maruti Suzuki	Swift	2017	7	83,500 km	Manual	second	
4	Maruti Suzuki	Baleno	2019	5	45,000 km	Automatic	first	

	FuelType	PostedDate	AskPrice
0	Petrol	Nov-24	1,95,000
1	Diesel	Jul-24	3,75,000
2	Diesel	Nov-24	1,84,999
3	Diesel	Nov-24	5,65,000
4	Petrol	Nov-24	6,85,000

```
[158]: data.shape
```

```
[158]: (9582, 10)
```

```
[159]: # Clean the AskPrice column
data['AskPrice'] = (
    data['AskPrice']
    .str.replace(' ', '', regex=False) # Remove rupee symbol
    .str.replace(',', '', regex=False) # Remove commas
    .str.strip() # Remove any leading/trailing spaces
    .astype(int) # Convert to integer
)
```

```
[160]: data.head()
```

```
[160]:
```

	Brand	model	Year	Age	kmDriven	Transmission	Owner	\
0	Honda	City	2001	23	98,000 km	Manual	second	
1	Toyota	Innova	2009	15	190000.0 km	Manual	second	
2	Volkswagen	VentoTest	2010	14	77,246 km	Manual	first	
3	Maruti Suzuki	Swift	2017	7	83,500 km	Manual	second	
4	Maruti Suzuki	Baleno	2019	5	45,000 km	Automatic	first	

	FuelType	PostedDate	AskPrice
0	Petrol	Nov-24	195000
1	Diesel	Jul-24	375000
2	Diesel	Nov-24	184999
3	Diesel	Nov-24	565000

4 Petrol Nov-24 685000

```
[161]: # Clean the kmDriven column
data['kmDriven'] = (
    data['kmDriven']
    .str.replace('km', '', regex=False) # Remove rupee symbol
    .str.replace(',', '', regex=False) # Remove commas
    .str.strip() # Remove any leading/trailing spaces
    .astype(float) # Convert to integer
)
```

```
[162]: data.head()
```

```
[162]:
```

	Brand	model	Year	Age	kmDriven	Transmission	Owner	\
0	Honda	City	2001	23	98000.0	Manual	second	
1	Toyota	Innova	2009	15	190000.0	Manual	second	
2	Volkswagen	VentoTest	2010	14	77246.0	Manual	first	
3	Maruti Suzuki	Swift	2017	7	83500.0	Manual	second	
4	Maruti Suzuki	Baleno	2019	5	45000.0	Automatic	first	

	FuelType	PostedDate	AskPrice
0	Petrol	Nov-24	195000
1	Diesel	Jul-24	375000
2	Diesel	Nov-24	184999
3	Diesel	Nov-24	565000
4	Petrol	Nov-24	685000

```
[173]: data['kmDriven'] = data['kmDriven'].fillna(data['kmDriven'].mean())
```

```
[174]: data.isna().sum()
```

```
[174]: Brand          0
model            0
Year            0
Age             0
kmDriven        0
Transmission    0
Owner           0
FuelType        0
PostedDate      0
AskPrice        0
dtype: int64
```

```
[175]: data.drop(columns = ['Age'],axis=1,inplace=True)
```

```
[176]: data.head()
```

```
[176]:
```

	Brand	model	Year	kmDriven	Transmission	Owner	FuelType	\
0	Honda	City	2001	98000.0	Manual	second	Petrol	
1	Toyota	Innova	2009	190000.0	Manual	second	Diesel	
2	Volkswagen	VentoTest	2010	77246.0	Manual	first	Diesel	
3	Maruti Suzuki	Swift	2017	83500.0	Manual	second	Diesel	
4	Maruti Suzuki	Baleno	2019	45000.0	Automatic	first	Petrol	

	PostedDate	AskPrice
0	Nov-24	195000
1	Jul-24	375000
2	Nov-24	184999
3	Nov-24	565000
4	Nov-24	685000

```
[177]: data['PostedDate'] = pd.to_datetime(data['PostedDate'] + '-01',
↳format='%b-%y-%d')
```

```
[178]: # Extract Year and Month from the PostedDate into separate columns
data['Year_Posted'] = data['PostedDate'].dt.year
data['Month_Posted'] = data['PostedDate'].dt.month
```

```
[179]: data.head()
```

```
[179]:
```

	Brand	model	Year	kmDriven	Transmission	Owner	FuelType	\
0	Honda	City	2001	98000.0	Manual	second	Petrol	
1	Toyota	Innova	2009	190000.0	Manual	second	Diesel	
2	Volkswagen	VentoTest	2010	77246.0	Manual	first	Diesel	
3	Maruti Suzuki	Swift	2017	83500.0	Manual	second	Diesel	
4	Maruti Suzuki	Baleno	2019	45000.0	Automatic	first	Petrol	

	PostedDate	AskPrice	Year_Posted	Month_Posted
0	2024-11-01	195000	2024	11
1	2024-07-01	375000	2024	7
2	2024-11-01	184999	2024	11
3	2024-11-01	565000	2024	11
4	2024-11-01	685000	2024	11

```
[180]: data.drop(columns = 'PostedDate',axis=1, inplace=True)
```

```
[181]: data.head()
```

```
[181]:
```

	Brand	model	Year	kmDriven	Transmission	Owner	FuelType	\
0	Honda	City	2001	98000.0	Manual	second	Petrol	
1	Toyota	Innova	2009	190000.0	Manual	second	Diesel	
2	Volkswagen	VentoTest	2010	77246.0	Manual	first	Diesel	
3	Maruti Suzuki	Swift	2017	83500.0	Manual	second	Diesel	
4	Maruti Suzuki	Baleno	2019	45000.0	Automatic	first	Petrol	

	AskPrice	Year_Posted	Month_Posted
0	195000	2024	11
1	375000	2024	7
2	184999	2024	11
3	565000	2024	11
4	685000	2024	11

```
[182]: data['Month_Posted'].value_counts()
```

```
[182]: 11    8693
      10    616
      9    145
      8     63
      7     29
      6     18
      4      6
      5      5
      12     3
      2      2
      1      1
      3      1
      Name: Month_Posted, dtype: int64
```

```
[183]: data['Year_Posted'].value_counts()
```

```
[183]: 2024    9579
      2023      3
      Name: Year_Posted, dtype: int64
```

```
[184]: # Replace 'Manual' with 0 and 'Automatic' with 1 in the 'Transmission' column
data['Transmission'] = data['Transmission'].replace({'Manual': 0, 'Automatic': 1})
```

```
[185]: data.head()
```

```
[185]:
```

	Brand	model	Year	kmDriven	Transmission	Owner	FuelType	\
0	Honda	City	2001	98000.0	0	second	Petrol	
1	Toyota	Innova	2009	190000.0	0	second	Diesel	
2	Volkswagen	VentoTest	2010	77246.0	0	first	Diesel	
3	Maruti Suzuki	Swift	2017	83500.0	0	second	Diesel	
4	Maruti Suzuki	Baleno	2019	45000.0	1	first	Petrol	

	AskPrice	Year_Posted	Month_Posted
0	195000	2024	11
1	375000	2024	7
2	184999	2024	11

```
3    565000      2024      11
4    685000      2024      11
```

```
[186]: data['FuelType'] = data['FuelType'].replace({'Diesel': 0, 'Petrol': 1, 'Hybrid/
↪CNG':2})
```

```
[187]: data.head()
```

```
[187]:
```

	Brand	model	Year	kmDriven	Transmission	Owner	FuelType	\
0	Honda	City	2001	98000.0	0	second	1	
1	Toyota	Innova	2009	190000.0	0	second	0	
2	Volkswagen	VentoTest	2010	77246.0	0	first	0	
3	Maruti Suzuki	Swift	2017	83500.0	0	second	0	
4	Maruti Suzuki	Baleno	2019	45000.0	1	first	1	

	AskPrice	Year_Posted	Month_Posted
0	195000	2024	11
1	375000	2024	7
2	184999	2024	11
3	565000	2024	11
4	685000	2024	11

```
[188]: data['Owner'].value_counts()
```

```
[188]: first      4800
second    4782
Name: Owner, dtype: int64
```

```
[189]: data['Owner'] = data['Owner'].replace({'first': 0, 'second': 1})
```

```
[190]: data.head()
```

```
[190]:
```

	Brand	model	Year	kmDriven	Transmission	Owner	FuelType	\
0	Honda	City	2001	98000.0	0	1	1	
1	Toyota	Innova	2009	190000.0	0	1	0	
2	Volkswagen	VentoTest	2010	77246.0	0	0	0	
3	Maruti Suzuki	Swift	2017	83500.0	0	1	0	
4	Maruti Suzuki	Baleno	2019	45000.0	1	0	1	

	AskPrice	Year_Posted	Month_Posted
0	195000	2024	11
1	375000	2024	7
2	184999	2024	11
3	565000	2024	11
4	685000	2024	11

```
[191]: data['model'].value_counts()
```



```
[191]: City          330
      Wagon-R        311
      Swift          283
      Creta          260
      Ertiga         249
      ...
      H5x            1
      Punch          1
      Tiguan All Space 1
      Cedia          1
      Gran Turismo   1
      Name: model, Length: 400, dtype: int64
```

```
[192]: from sklearn.preprocessing import LabelEncoder
      # Initialize the LabelEncoder
      le = LabelEncoder()
```

```
[193]: # Convert the 'Model' column to numeric values
      data['model'] = le.fit_transform(data['model'])
```

```
[194]: data.head()
```

```
[194]:      Brand  model  Year  kmDriven  Transmission  Owner  FuelType  \
0      Honda    84  2001   98000.0             0      1          1
1      Toyota   187  2009  190000.0             0      1          0
2  Volkswagen   347  2010   77246.0             0      0          0
3  Maruti Suzuki  317  2017   83500.0             0      1          0
4  Maruti Suzuki    52  2019   45000.0             1      0          1

      AskPrice  Year_Posted  Month_Posted
0    195000      2024         11
1    375000      2024          7
2    184999      2024         11
3    565000      2024         11
4    685000      2024         11
```

```
[195]: data['Brand'] = le.fit_transform(data['Brand'])
```

```
[196]: data.head()
```

```
[196]:      Brand  model  Year  kmDriven  Transmission  Owner  FuelType  AskPrice  \
0      12     84  2001   98000.0             0      1          1    195000
1      36    187  2009  190000.0             0      1          0    375000
2      37    347  2010   77246.0             0      0          0    184999
3      23    317  2017   83500.0             0      1          0    565000
4      23     52  2019   45000.0             1      0          1    685000
```

	Year_Posted	Month_Posted
0	2024	11
1	2024	7
2	2024	11
3	2024	11
4	2024	11

```
[197]: # split the data into X and Y
x= data.drop(columns = 'AskPrice')
y = data['AskPrice']
```

```
[198]: x.head()
```

```
[198]:
```

	Brand	model	Year	kmDriven	Transmission	Owner	FuelType	Year_Posted	\
0	12	84	2001	98000.0	0	1	1	2024	
1	36	187	2009	190000.0	0	1	0	2024	
2	37	347	2010	77246.0	0	0	0	2024	
3	23	317	2017	83500.0	0	1	0	2024	
4	23	52	2019	45000.0	1	0	1	2024	

	Month_Posted
0	11
1	7
2	11
3	11
4	11

```
[199]: y.head()
```

```
[199]:
```

0	195000
1	375000
2	184999
3	565000
4	685000

Name: AskPrice, dtype: int32

```
[200]: # normalization

from sklearn.preprocessing import MinMaxScaler
# Initialize MinMaxScaler
scaler = MinMaxScaler()
```

```
[201]: data_normalized = pd.DataFrame(scaler.fit_transform(data), columns=data.columns)
```

```
[202]: data_normalized.head()
```

```
[202]:
```

	Brand	model	Year	kmDriven	Transmission	Owner	FuelType	\
0	0.315789	0.210526	0.394737	0.100000	0.0	1.0	0.5	
1	0.947368	0.468672	0.605263	0.193877	0.0	1.0	0.0	
2	0.973684	0.869674	0.631579	0.078822	0.0	0.0	0.0	
3	0.605263	0.794486	0.815789	0.085204	0.0	1.0	0.0	
4	0.605263	0.130326	0.868421	0.045918	1.0	0.0	0.5	

	AskPrice	Year_Posted	Month_Posted
0	0.004237	1.0	0.909091
1	0.008474	1.0	0.545455
2	0.004001	1.0	0.909091
3	0.012946	1.0	0.909091
4	0.015770	1.0	0.909091

```
[203]: from sklearn.model_selection import train_test_split
```

```
[204]: X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
↳random_state=42)
```

```
[205]: # Display the shapes of the resulting datasets
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: (7665, 9)
X_test shape: (1917, 9)
y_train shape: (7665,)
y_test shape: (1917,)
```

```
[206]: #building the XgBoost model
```

```
!pip install xgboost
```

```
Requirement already satisfied: xgboost in c:\users\user\anaconda3\lib\site-
packages (2.1.2)
Requirement already satisfied: numpy in c:\users\user\anaconda3\lib\site-
packages (from xgboost) (1.24.4)
Requirement already satisfied: scipy in c:\users\user\anaconda3\lib\site-
packages (from xgboost) (1.10.1)
```

```
[207]: import xgboost as xgb
from sklearn.metrics import mean_squared_error, r2_score
```

```
[208]: # Initialize the XGBoost regressor model
xg_reg = xgb.XGBRegressor(objective='reg:squarederror', eval_metric='rmse',
↳random_state=42)
```

```
# Train the model
xg_reg.fit(X_train, y_train)

# Predict on the test set
y_pred = xg_reg.predict(X_test)
```

```
[209]: y_pred
```

```
[209]: array([ 346213.9 , 302984. , 1098718.4 , ..., 661576.56, 108354.48,
        685818.9 ], dtype=float32)
```

```
[210]: # Evaluate the model
rmse = mean_squared_error(y_test, y_pred, squared=False)
r2 = r2_score(y_test, y_pred)

# Output the results
print(f'Root Mean Squared Error (RMSE): {rmse}')
print(f'R2 Score: {r2}')
```

```
Root Mean Squared Error (RMSE): 769862.3129815074
R2 Score: 0.7788931419558567
```

```
[ ]:
```

```
[211]: from sklearn.ensemble import AdaBoostRegressor
from sklearn.tree import DecisionTreeRegressor
```

```
[212]: # Initialize the AdaBoost Regressor with a weak learner (Decision Tree)
adaboost_regressor =
    ↳AdaBoostRegressor(base_estimator=DecisionTreeRegressor(max_depth=2),
    ↳n_estimators=50, random_state=42)

# Train the model
adaboost_regressor.fit(X_train, y_train)
```

```
C:\Users\USER\anaconda3\Lib\site-packages\sklearn\ensemble\_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version 1.2 and
will be removed in 1.4.
    warnings.warn(
```

```
[212]: AdaBoostRegressor(base_estimator=DecisionTreeRegressor(max_depth=2),
                        random_state=42)
```

```
[213]: # Predict on the test set
y_pred1 = adaboost_regressor.predict(X_test)
```

```
[214]: y_pred1
```

```
[214]: array([1062533.33455545, 1062533.33455545, 1386868.25885958, ...,  
        1386868.25885958, 1062533.33455545, 1062533.33455545])
```

```
[215]: # Evaluate the model  
rmse = mean_squared_error(y_test, y_pred, squared=False) # Root Mean Squared  
↳Error  
r2 = r2_score(y_test, y_pred) # R2 Score  
  
# Print results  
print(f'Root Mean Squared Error (RMSE): {rmse}')
```

```
print(f'R2 Score: {r2}')
```

Root Mean Squared Error (RMSE): 769862.3129815074
R2 Score: 0.7788931419558567

```
[ ]:
```

```
[ ]:
```